Astra - AI Paid Features Spec (v1)

For the development session working on github.com/Aykhan-iskandarli/astro. Scope: the six AI-powered paid/funnel features below. **Payment provider and checkout are OUT OF SCOPE** (decided separately) - build the AI layer so that access can later be gated through the existing `AuthProtectedComponent` hook and the planned `ai.interpretation` module described in `docs/architecture.md`.

1. Global rules (apply to every feature)

- **Server-side only.** Claude API is called from `apps/api`; the API key comes from an environment variable and never reaches the frontend or the repo.
- **Grounded in engine data.** The model receives the JSON returned by the existing engine endpoints and only interprets it. It must never calculate or invent positions, degrees, houses, aspects or dates. The system prompt must state this, and outputs must only reference data present in the input.
- **Languages.** Output language = UI locale (az / en). Azerbaijani quality must be reviewed by a native speaker before launch.
- **Structured output.** The model returns JSON (sections with title + text). The frontend renders it in the site's existing visual style; the same JSON feeds the PDF export.
- **Caching.** Cache results by hash of (input JSON + product + locale + prompt version). The same chart never pays twice.
- **Cost control.** Rate limit per IP/client, a global daily spend cap (stop AI calls when reached), and token-usage logging. Never log birth data.
- **Prompt versioning.** Prompts live in versioned files; the version is part of the cache key.
- **Safety / tone.** Entertainment disclaimer on every AI output. No medical, legal or financial instructions, no death or illness predictions, no fatalistic certainty; phrase as tendencies.
- **Model choice.** Before building features, run a test: 5 sample outputs per feature with Claude Haiku 5.5 and Claude Sonnet 5.5, in az and en, reviewed by a human. Prices seen in secondary sources (verify on the official Claude pricing page): Haiku 5.5 $0.10 input / $0.50 output per 1M tokens (prompts up to 100K), Sonnet 5.5 $2 / $10. Initial assumption below: Sonnet for paid reports, Haiku for teaser and weekly texts - confirm with the test.

2. Feature 1 - Full Personality Report (one-time purchase) - MAIN PRODUCT

Data	`POST /api/charts/natal`, `/destiny-matrix`, `/numerology` results for the same person.
Output	One coherent report synthesising all three systems (not a list of per-placement snippets): 1) Core personality, 2) Emotional world, 3) Career and talents, 4) Money and values, 5) Love and relationships, 6) Strengths and challenges, 7) Life lessons and growth path, 8) Short summary. Each section must cite which placements it is based on.
Model	Sonnet 5.5 (initial assumption).
Delivery	Rendered page + PDF download.

3. Feature 2 - Compatibility Report (one-time purchase)

Data	`POST /api/charts/synastry` (including the score) and `/composite` for two people.
Output	Relationship dynamics, attraction points, friction points and how to handle them, communication style, long-term potential, practical advice. Must not declare a relationship 'doomed' or 'perfect'.
Model	Sonnet 5.5 (initial assumption).
Delivery	Page + PDF; giftable (can be bought for a couple).

4. Feature 3 - Annual Forecast (one-time purchase)

Data	`POST /api/charts/timeline` (12 months), `/returns` (solar return), `/progressions`.
Output	Year theme overview, then month by month: main theme, key dates (only dates present in the timeline data), opportunities, cautions, advice. Must build on the existing free transit calendar, not duplicate it.
Model	Sonnet 5.5 (initial assumption). Longest output - expect roughly 2x the cost of Feature 1.
Delivery	Page + PDF.

5. Feature 4 - 'Ask Your Chart' Chat (question packs)

Data	The user's natal chart JSON (optionally transits for the current date) placed in the system prompt.
Output	Answers to free-form questions, grounded only in the chart data; politely refuses medical/legal/financial decisions and off-topic requests.
Cost	Use prompt caching for the chart JSON so follow-up questions are cheap. Enforce max questions per pack and max conversation length.
Note	Needs a way to track remaining questions per buyer - design the data model, leave payment integration out.

6. Feature 5 - Weekly Personal Forecast (subscription)

Data	`POST /api/charts/transits` for the coming week against the subscriber's natal chart.
Output	Short weekly text (about 150-250 words): week theme, 2-3 notable transits with dates, one piece of advice.
Model	Haiku 5.5 (initial assumption) - very low cost per message.
Note	Requires storing birth data + email server-side, which the app does not do today. Must be explicit opt-in with consent, covered in the privacy policy, with delete-my-data support.

7. Feature 6 - Free AI Teaser (conversion funnel)

Data	Same input as Feature 1.
Output	Only the opening paragraph of the Full Personality Report, shown free under the Natal chart tab, followed by a 'Read the full report' call to action.
Model	Haiku 5.5 (initial assumption); cached per chart; strict rate limit since it is free.
Goal	Show the quality of the paid report so users convert.

8. Rollout order

- Step 0: model test (Haiku vs Sonnet, az/en) - report results and wait for owner approval.
- Then: Feature 1 -> Feature 6 -> Feature 2 -> Feature 3 -> Feature 4 -> Feature 5.
- After each feature: short report (what was built, what was tested, measured tokens/cost per output, known gaps) and wait for approval.

9. Out of scope for now

- Payment provider selection and checkout integration.
- Pricing of each product (decided by the owner).
- Full user accounts beyond what Feature 4 and Feature 5 minimally need (design only, confirm with owner first).